

# CrowdROM: Shift-Aligned POD–MVAR Pipeline

Algorithms for Density-Field Reduced-Order Modeling  
with Hellinger Embedding and Simplex Mass Correction

Macias-Olivares, E.

February 2026

## 1 Problem Setting

We consider  $N$  interacting agents on a periodic domain  $\Omega = [0, L_x) \times [0, L_y)$  evolving under a discrete-time Vicsek–Morse model (??). At each time  $t_k = k \Delta t$ , the particle positions are  $\{\mathbf{x}_i^k\}_{i=1}^N \subset \Omega$ .

We construct a *density representation* by binning and smoothing:

$$\rho^k(x, y) = (G_\sigma * H^k)(x, y), \quad H_{jl}^k = \frac{1}{\Delta x \Delta y} \sum_{i=1}^N \mathbf{1}[\mathbf{x}_i^k \in C_{jl}], \quad (1)$$

where  $H^k \in \mathbb{R}^{N_y \times N_x}$  is the histogram on cells  $C_{jl} = [x_j, x_{j+1}) \times [y_l, y_{l+1})$  with spacing  $\Delta x = L_x/N_x$ ,  $\Delta y = L_y/N_y$ , and  $G_\sigma$  is a 2D Gaussian kernel with bandwidth  $\sigma$  (in grid-cell units), applied with *periodic wrapping* (?). The density satisfies  $\sum_{jl} \rho_{jl}^k \Delta x \Delta y = N$ .

In code, this is `density.compute_density_grid()` using `scipy.ndimage.gaussian_filter` with `mode='wrap'`.

We flatten each density frame to a column vector  $\boldsymbol{\rho}^k \in \mathbb{R}^n$ ,  $n = N_x N_y$ , and collect  $M$  training trajectories of  $T$  subsampled frames each into the *snapshot matrix*  $\mathbf{X}_{\text{raw}} = [\boldsymbol{\rho}_1^1, \dots, \boldsymbol{\rho}_1^T, \dots, \boldsymbol{\rho}_M^1, \dots, \boldsymbol{\rho}_M^T]^\top \in \mathbb{R}^{MT \times n}$ .

## 2 Step 1: Shift Alignment (sPOD)

Particle systems with net translational motion produce density fields that *drift* across  $\Omega$ . Standard POD is inefficient for traveling structures: a simple translation requires many Fourier-type modes to represent (??).

**Shifted POD (sPOD)** factors out integer-pixel translations *before* computing the POD basis, so that the reduced coordinates only need to capture structural deformation.

### 2.1 FFT Cross-Correlation Registration

Given a reference field  $\rho_{\text{ref}} \in \mathbb{R}^{N_y \times N_x}$  and a target frame  $\rho^k$ , we seek the cyclic shift  $(\delta_y^k, \delta_x^k) \in \mathbb{Z}^2$  that maximises the cross-correlation:

$$(\delta_y^k, \delta_x^k) = \arg \max_{(s_y, s_x) \in \mathbb{Z}_{N_y} \times \mathbb{Z}_{N_x}} \sum_{j,l} \rho_{\text{ref}}(j, l) \rho^k((j + s_y) \bmod N_y, (l + s_x) \bmod N_x). \quad (2)$$

---

**Algorithm 1** Shift Alignment (sPOD) — Training

---

**Require:** Snapshot matrix  $\mathbf{X}_{\text{raw}} \in \mathbb{R}^{MT \times n}$  reshaped as  $\{\rho^k\}_{k=1}^{MT}$ , each  $\rho^k \in \mathbb{R}^{N_y \times N_x}$

**Ensure:** Aligned snapshots  $\{\tilde{\rho}^k\}$ , reference  $\rho_{\text{ref}}$ , shift array  $\Delta \in \mathbb{Z}^{MT \times 2}$

```

1:  $\rho_{\text{ref}} \leftarrow \frac{1}{MT} \sum_{k=1}^{MT} \rho^k$                                      ▷ shift_align_ref='mean'
2: for  $k = 1, \dots, MT$  do
3:    $\hat{F}_{\text{ref}} \leftarrow \text{fft2}(\rho_{\text{ref}})$ ;  $\hat{F}_k \leftarrow \text{fft2}(\rho^k)$ 
4:    $C \leftarrow \text{Re}(\text{ifft2}(\hat{F}_{\text{ref}} \odot \hat{F}_k))$ 
5:    $(\delta_y^k, \delta_x^k) \leftarrow \arg \max_{(s_y, s_x)} C_{s_y, s_x}$                                 ▷ wrap to signed
6:    $\tilde{\rho}^k \leftarrow \text{roll}(\rho^k, \delta_y^k, \delta_x^k)$                                          ▷ Eq. (??)
7: end for
8: return  $\{\tilde{\rho}^k\}, \rho_{\text{ref}}, \Delta$ 

```

---

**Algorithm 2** Shift Alignment — Test-Time Un-alignment

---

**Require:** Test density  $\rho_{\text{test}}^{1:T_{\text{test}}}$ , training reference  $\rho_{\text{ref}}$ , forecast start index  $T_0$

**Ensure:** Aligned test density, predicted shifts for forecast

```

1: for  $k = 1, \dots, T_{\text{test}}$  do
2:    $(\delta_y^k, \delta_x^k) \leftarrow \text{fft\_cross\_correlation\_shift}(\rho_{\text{ref}}, \rho_{\text{test}}^k)$ 
3:    $\tilde{\rho}_{\text{test}}^k \leftarrow \text{roll}(\rho_{\text{test}}^k, \delta_y^k, \delta_x^k)$ 
4: end for
5: Fit  $\hat{\delta}(t) = \mathbf{c}_0 + \mathbf{c}_1 t$  to  $\{\delta^k\}_{k=1}^{T_0}$  via least-squares                               ▷ predict_shifts_linear
6: Extrapolate  $\hat{\delta}^k$  for  $k = T_0 + 1, \dots, T_0 + H$ 
7: Un-align predictions:  $\hat{\rho}_{\text{phys}}^k \leftarrow \text{roll}(\tilde{\rho}^k, -\hat{\delta}_y^k, -\hat{\delta}_x^k)$ 

```

---

By the convolution theorem on the periodic domain, this is computed in  $O(n \log n)$  via the 2D FFT:

$$C = \mathcal{F}^{-1}[\mathcal{F}[\rho_{\text{ref}}] \odot \overline{\mathcal{F}[\rho^k]}], \quad (\delta_y^k, \delta_x^k) = \arg \max_{(s_y, s_x)} \text{Re}(C_{s_y, s_x}), \quad (3)$$

where  $(\cdot)$  denotes elementwise complex conjugation and  $\odot$  is the Hadamard product. The shift is converted to signed integers by wrapping indices exceeding  $N_y/2$  or  $N_x/2$ .

## 2.2 Alignment and Un-alignment

The aligned frame is obtained by cyclic rolling:

$$\tilde{\rho}^k(j, l) = \rho^k((j - \delta_y^k) \bmod N_y, (l - \delta_x^k) \bmod N_x). \quad (4)$$

At prediction time, forecasted (aligned) fields are *un-aligned* by applying the *negative* shift:

$$\hat{\rho}_{\text{phys}}^k(j, l) = \tilde{\rho}^k((j + \delta_y^k) \bmod N_y, (l + \delta_x^k) \bmod N_x). \quad (5)$$

For forecast frames beyond the teacher-forced window, shifts are *linearly extrapolated*:  $\hat{\delta}(t) = \mathbf{c}_0 + \mathbf{c}_1 t$ , fit by least-squares to the known shifts.

**Implementation.** Defined in `src/rectsim/shift_align.py`: `fft_cross_correlation_shift()`, `compute_reference_field()`, `align_training_data()`, `apply_shifts()`/`undo_shifts()`, `predict_shifts_linear()`. Called from `pod_builder.build_pod_basis()` (training) and `test_evaluator.evaluate_test_runs()` (testing). Config keys: `rom.shift_align: true`, `rom.shift_align_ref: "mean"`.

### 3 Step 2: Hellinger (Square-Root) Density Transform

After alignment, we apply a variance-stabilising transform before POD. The *square-root transform* embeds densities into the *Hellinger space* (??):

$$f(\rho) = \sqrt{\rho + \varepsilon}, \quad \varepsilon = 10^{-10}, \quad (6)$$

applied element-wise. The small  $\varepsilon$  prevents division issues at zero-density cells.

#### 3.1 Why It Helps POD

1. **Dynamic range compression.** Raw density fields are non-negative and spiky: a few bright cluster peaks dominate the  $L^2$  variance. The square root compresses peaks ( $\sqrt{\rho_{\max}}$ ) and lifts tails ( $\sqrt{\rho_{\min}}$ ), distributing energy across more POD modes and requiring fewer modes for a given energy threshold.
2. **Hellinger metric.** In the transformed space, the  $L^2$  distance equals the Hellinger distance (?):
 
$$\|\sqrt{\rho_1} - \sqrt{\rho_2}\|_2^2 = H^2(\rho_1, \rho_2) = \frac{1}{2} \int (\sqrt{\rho_1} - \sqrt{\rho_2})^2, \quad (7)$$

which is a proper statistical divergence between density functions, bounded in [0, 1] for probability measures. Minimising POD reconstruction error in Hellinger space is therefore statistically principled.

3. **Variance stabilisation.** For Poisson-like count data (as in histogrammed particles),  $\text{Var}[\rho] \propto \rho$ . The Anscombe transform  $\sqrt{\rho + 3/8}$  is the classical variance-stabiliser (?). Our  $\sqrt{\rho + \varepsilon}$  serves the same role, making the POD residual approximately homoscedastic.

#### 3.2 Inverse Transform

After POD lifting, the reconstructed field  $\hat{\mathbf{y}}$  lies in the *transformed* space. The inverse is:

$$\hat{\rho} = [\max(\hat{\mathbf{y}}, 0)]^2 - \varepsilon. \quad (8)$$

The  $\max(\cdot, 0)$  clamp is necessary because POD reconstruction (a linear projection) can produce negative values in the transformed space, even though true  $\sqrt{\rho + \varepsilon} > 0$ .

**Implementation.** **Forward:** `pod_builder.py` lines 118–120: `X_all = np.sqrt(X_all + density_transform_0)`. Applied *after* shift alignment, *before* POD mean-centering. **Inverse:** `test_evaluator.py` lines 355–357: `np.maximum(pred, 0.0)**2 - density_transform_eps`. Config key: `rom.density_transform: "sqrt"`, `rom.density_transform_eps: 1e-10`.

## 4 Step 3: Proper Orthogonal Decomposition (POD)

Given the aligned, transformed snapshot matrix  $\tilde{\mathbf{X}} \in \mathbb{R}^{MT \times n}$ :

1. **Mean-center:**  $\boldsymbol{\mu} = \frac{1}{MT} \sum_k \tilde{\mathbf{x}}_k$ ,  $\bar{\mathbf{X}} = \tilde{\mathbf{X}} - \mathbf{1}\boldsymbol{\mu}^\top$ .
2. **SVD:**  $\bar{\mathbf{X}}^\top = \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^\top$ ,  $\mathbf{U} \in \mathbb{R}^{n \times MT}$ ,  $\sigma_1 \geq \sigma_2 \geq \dots \geq 0$ .
3. **Truncate** to  $d$  modes, determined by either:
  - *Fixed:*  $d = \text{rom.fixed\_modes}$  (e.g.  $d = 19$ ), or
  - *Energy threshold:* smallest  $d$  such that  $\sum_{i=1}^d \sigma_i^2 / \sum_{i=1}^{\min(MT,n)} \sigma_i^2 \geq E_{\text{target}}$  (e.g.  $E_{\text{target}} = 0.99$ ).

$$\mathbf{U}_d = [\mathbf{u}_1, \dots, \mathbf{u}_d] \in \mathbb{R}^{n \times d}.$$

4. **Project to latent space:**  $\mathbf{a}^k = \mathbf{U}_d^\top (\tilde{\mathbf{x}}_k - \boldsymbol{\mu}) \in \mathbb{R}^d$ .
5. **Lift (reconstruct):**  $\hat{\mathbf{x}}_k = \mathbf{U}_d \hat{\mathbf{a}}^k + \boldsymbol{\mu}$ .

**Implementation.** `pod_builder.build_pod_basis()`: uses `np.linalg.svd(X_centered.T, full_matrices=False)`. Fixed modes from `rom.fixed_modes`, energy from `rom.pod_energy`.

## 5 Step 4: MVAR Latent Dynamics

In the latent space  $\{\mathbf{a}^k \in \mathbb{R}^d\}$ , we model the dynamics with a *Vector AutoRegressive model of order  $p$*  (VAR( $p$ )) with ridge regularisation (??):

$$\mathbf{a}^{k+1} = \sum_{j=1}^p \mathbf{A}_j \mathbf{a}^{k+1-j} + \mathbf{b} + \boldsymbol{\epsilon}^k, \quad \boldsymbol{\epsilon}^k \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma}_\epsilon), \quad (9)$$

where  $\mathbf{A}_j \in \mathbb{R}^{d \times d}$  are the coefficient matrices,  $\mathbf{b} \in \mathbb{R}^d$  is the intercept, and  $p$  is the lag order.

### 5.1 Training

Stack the lagged history into a design row:

$$\mathbf{x}_{\text{train}}^k = [\mathbf{a}^{k-p}, \mathbf{a}^{k-p+1}, \dots, \mathbf{a}^{k-1}]^\top \in \mathbb{R}^{pd}, \quad \mathbf{y}_{\text{train}}^k = \mathbf{a}^k \in \mathbb{R}^d. \quad (10)$$

The coefficient matrix  $\mathbf{W} = [\mathbf{A}_p, \dots, \mathbf{A}_1, \mathbf{b}] \in \mathbb{R}^{d \times (pd+1)}$  is found by ridge regression:

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \sum_k \|\mathbf{y}^k - \mathbf{W}[\mathbf{x}^k; 1]\|_2^2 + \alpha \|\mathbf{W}\|_F^2, \quad (11)$$

with regularisation  $\alpha > 0$  (config: `rom.models.mvar.ridge_alpha`). This has the closed-form solution via ridge normal equations, implemented using `sklearn.linear_model.Ridge(fit_intercept=True)`.

---

**Algorithm 3** Euclidean Projection onto the Scaled Simplex  $\mathcal{S}_{M_0}$ 


---

**Require:**  $\hat{\rho} \in \mathbb{R}^n$ , mass target  $M_0 > 0$

**Ensure:**  $\rho^* = \Pi_{\mathcal{S}_{M_0}}(\hat{\rho})$

- 1:  $\mathbf{v} \leftarrow \hat{\rho}/M_0$  ▷ scale to unit simplex
  - 2: Sort:  $u_1 \geq u_2 \geq \dots \geq u_n$  (descending)
  - 3:  $S_j \leftarrow \sum_{i=1}^j u_i - 1$ ,  $j = 1, \dots, n$  ▷ running cumulative sum
  - 4:  $\hat{j} \leftarrow \max\{j \in [n] \mid u_j - S_j/j > 0\}$
  - 5:  $\theta \leftarrow S_{\hat{j}}/\hat{j}$  ▷ dual variable / Lagrange multiplier
  - 6:  $\rho_i^* \leftarrow \max(v_i - \theta, 0) \cdot M_0$ ,  $i = 1, \dots, n$
  - 7: **return**  $\rho^*$
- 

## 5.2 Autoregressive Forecasting

Given the last  $p$  known latent states  $\mathbf{a}^{T_0-p+1}, \dots, \mathbf{a}^{T_0}$  (the “initial condition window”), the forecast is generated autoregressively:

$$\hat{\mathbf{a}}^{T_0+h} = \sum_{j=1}^p \mathbf{A}_j \hat{\mathbf{a}}^{T_0+h-j} + \mathbf{b}, \quad h = 1, 2, \dots, H, \quad (12)$$

where  $\hat{\mathbf{a}}^k = \mathbf{a}^k$  for  $k \leq T_0$  (teacher-forced prefix).

**Implementation.** `src/rectsim/mvar_trainer.py`: `train_mvar_model()`. Config: `rom.models.mvar.{lag, ridge_alpha}`.

## 6 Step 5: Simplex Mass Projection

The nonlinear inverse transform ( $\hat{\rho} = \max(\hat{y}, 0)^2 - \varepsilon$ ) does *not* preserve total mass:  $\sum_i \hat{\rho}_i \neq \sum_i \rho_i^0 = M_0$  because the square of a linear reconstruction is not linear (Jensen’s inequality). We correct this by an  $L^2$ -optimal projection onto the *scaled probability simplex* (?):

$$\mathcal{S}_{M_0} = \{\rho \in \mathbb{R}^n \mid \rho_i \geq 0 \forall i, \quad \sum_{i=1}^n \rho_i = M_0\}, \quad (13)$$

where  $M_0 = \sum_i \rho_i^{T_0}$  is the ground-truth total mass at the forecast start.

$$\Pi_{\mathcal{S}_{M_0}}(\hat{\rho}) = \arg \min_{\rho \in \mathcal{S}_{M_0}} \|\rho - \hat{\rho}\|_2^2. \quad (14)$$

### 6.1 Algorithm (Duchi et al., 2008)

**Complexity.**  $O(n \log n)$  due to the sort; all other operations are  $O(n)$ .

**KKT interpretation.** The Lagrangian of (??) on the unit simplex is  $L(\rho, \lambda, \mu) = \frac{1}{2} \|\rho - \hat{\rho}\|^2 + \lambda(\mathbf{1}^\top \rho - M_0) - \mu^\top \rho$ . The KKT conditions yield  $\rho_i^* = \max(\hat{\rho}_i - \lambda^*, 0)$  with  $\lambda^* = \theta \cdot M_0$  chosen so that  $\sum_i \rho_i^* = M_0$ . When  $\hat{\rho} \geq 0$  already (as after the  $\max(\cdot, 0)^2$  step), the projection reduces to a *uniform additive shift*  $\rho_i^* = \hat{\rho}_i - \lambda^*$  with truncation at zero, identical to the offset-then-clamp method.

**Implementation.** `test_evaluator._project_simplex(rho_frame, mass_target)`. Applied per frame after the inverse  $\sqrt{\rho}$  transform. Config: `rom.mass_postprocess: "simplex"`.

## 7 Complete Pipeline

### 7.1 Training Pipeline

1. **Simulate**  $M$  training trajectories (Vicsek–Morse,  $N$  agents,  $T_{\text{sim}}$  seconds). `cli.py`
2. **KDE**: bin particles onto  $N_y \times N_x$  grid, Gaussian-smooth with bandwidth  $\sigma$ . `density.py`

$$\boldsymbol{\rho}^k \in \mathbb{R}^{N_y \times N_x}, \quad \sum_{j,l} \rho_{jl} \Delta x \Delta y = N. \quad (15)$$

3. **Shift Alignment** (Alg. ??): compute reference  $\rho_{\text{ref}}$  and per-frame shifts  $\delta^k$ ; roll all frames to remove translational drift. `shift_align.py`

$$\tilde{\boldsymbol{\rho}}^k = \mathcal{T}_{\delta^k}[\boldsymbol{\rho}^k]. \quad (16)$$

4.  $\sqrt{\rho}$  **Transform**: embed aligned densities into Hellinger space. `pod_builder.py`

$$\tilde{\mathbf{x}}^k = \sqrt{\tilde{\boldsymbol{\rho}}^k + \varepsilon}. \quad (17)$$

5. **POD (SVD)**: mean-centre, compute SVD, truncate to  $d$  modes. `pod_builder.py`

$$\boldsymbol{\mu} = \bar{\mathbf{x}}, \quad (\tilde{\mathbf{X}} - \mathbf{1}\boldsymbol{\mu}^\top)^\top = \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^\top, \quad \mathbf{a}^k = \mathbf{U}_d^\top(\tilde{\mathbf{x}}^k - \boldsymbol{\mu}). \quad (18)$$

6. **MVAR**: fit VAR( $p$ ) with ridge  $\alpha$  on latent trajectories  $\{\mathbf{a}^k\}$ . `mvar_trainer.py`

$$\mathbf{W}^* = \arg \min_k \sum \|\mathbf{a}^{k+1} - \mathbf{W}[\mathbf{a}^{k-p+1:k}; 1]\|^2 + \alpha \|\mathbf{W}\|_F^2. \quad (19)$$

### 7.2 Prediction Pipeline (Test Time)

Given a test trajectory with density known up to  $t = t_0$  (the “conditioning window”):

1. **Shift-align** test density to training reference:  $\tilde{\rho}_{\text{test}}^k$ . Record shifts  $\delta_{\text{test}}^{1:T_0}$ . `test_evaluator.py`
2.  $\sqrt{\rho}$  **forward**:  $\tilde{x}_{\text{test}}^k = \sqrt{\tilde{\rho}_{\text{test}}^k + \varepsilon}$ .
3. **Project to latent**:  $\mathbf{a}_{\text{test}}^k = \mathbf{U}_d^\top(\tilde{x}_{\text{test}}^k - \boldsymbol{\mu})$ .
4. **MVAR forecast**: autoregressively generate  $\hat{\mathbf{a}}^{T_0+1}, \dots, \hat{\mathbf{a}}^{T_0+H}$ .
5. **Lift (POD inverse)**:  $\hat{\mathbf{y}}^k = \mathbf{U}_d \hat{\mathbf{a}}^k + \boldsymbol{\mu}$ .
6. **Inverse  $\sqrt{\rho}$**  (Eq. ??):  $\hat{\boldsymbol{\rho}}^k = [\max(\hat{\mathbf{y}}^k, 0)]^2 - \varepsilon$ .
7. **Simplex projection** (Alg. ??):  $\hat{\boldsymbol{\rho}}^k \leftarrow \Pi_{\mathcal{S}_{M_0}}(\hat{\boldsymbol{\rho}}^k)$ .
8. **Un-align shift**: extrapolate shifts for forecast horizon, reverse the cyclic roll.  $\hat{\boldsymbol{\rho}}_{\text{phys}}^k = \mathcal{T}_{-\hat{\delta}^k}[\hat{\boldsymbol{\rho}}^k]$ .

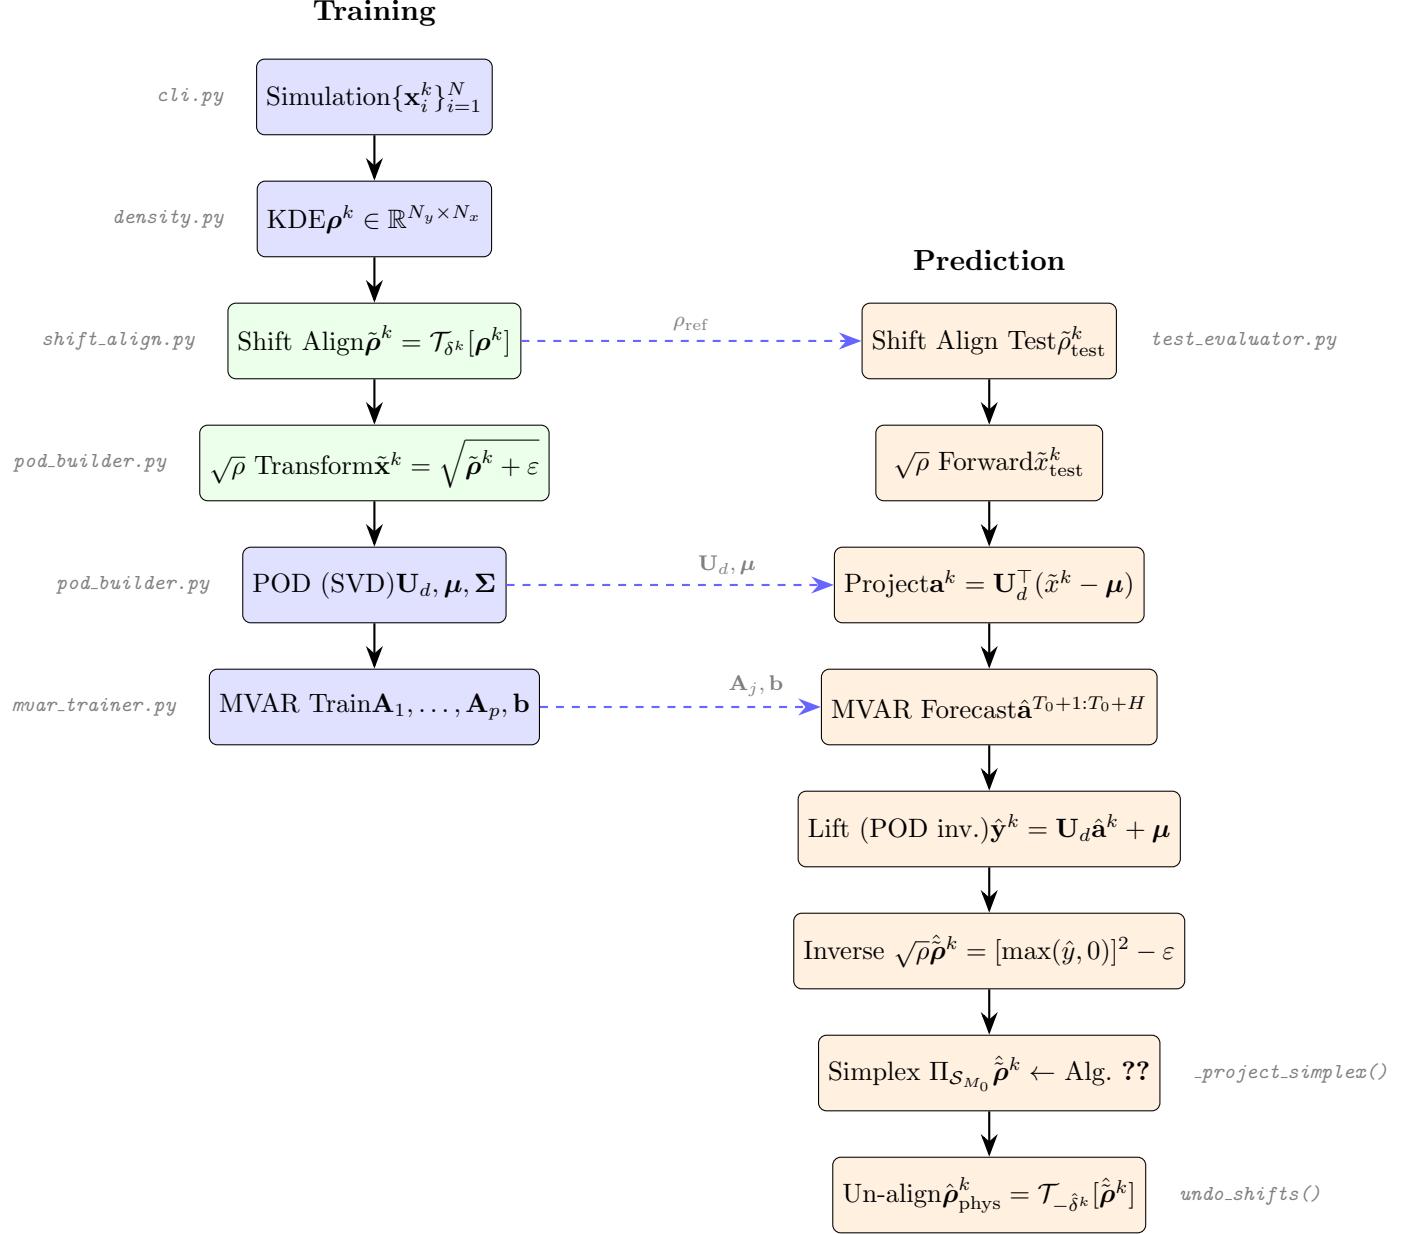


Figure 1: Complete CrowdROM pipeline. **Blue** = training only, **orange** = test time only, **green** = shared operations. Dashed arrows show training artifacts consumed at test time.

### 7.3 Pipeline Diagram

## 8 Configuration Summary

All pipeline options are set in the YAML config under the `rom:` block:

Config Key	Values	Effect
<code>shift_align</code>	true/false	Enable sPOD alignment
<code>shift_align_ref</code>	"mean", "first", "median"	Reference field computation
<code>density_transform</code>	"raw", "sqrt", "log"	Pre-POD transform
<code>density_transform_eps</code>	float ( $10^{-10}$ )	Regularisation $\varepsilon$
<code>fixed_modes</code>	int or null	Fixed POD rank $d$
<code>pod_energy</code>	float (0.99)	Energy threshold if no fixed rank
<code>mass_postprocess</code>	"none", "simplex", "scale"	Post-inverse mass correction
<code>models.mvar.lag</code>	int (5)	VAR order $p$
<code>models.mvar.ridge_alpha</code>	float ( $10^{-4}$ )	Ridge regularisation $\alpha$

Table 1: Pipeline configuration keys (`rom:` section of YAML).

## References

- Anscombe, F. J. (1948). The transformation of Poisson, binomial and negative-binomial data. *Biometrika*, 35(3/4):246–254.
- Black, F., Schulze, P., and Reiss, J. (2020). Projection-based model reduction with dynamically transformed modes. *ESAIM: Mathematical Modelling and Numerical Analysis*, 54(6):2011–2043.
- D’Orsogna, M. R., Chuang, Y.-L., Bertozzi, A. L., and Chayes, L. S. (2006). Self-propelled particles with soft-core interactions: patterns, stability, and collapse. *Physical Review Letters*, 96(10):104302.
- Duchi, J., Shalev-Shwartz, S., Singer, Y., and Chandra, T. (2008). Efficient projections onto the  $\ell_1$ -ball for learning in high dimensions. In *Proceedings of the 25th International Conference on Machine Learning (ICML)*, pages 272–279.
- Hastie, T., Tibshirani, R., and Friedman, J. (2009). *The Elements of Statistical Learning*. Springer, 2nd edition.
- Le Cam, L. (1986). *Asymptotic Methods in Statistical Decision Theory*. Springer-Verlag, New York.
- Lütkepohl, H. (2005). *New Introduction to Multiple Time Series Analysis*. Springer-Verlag, Berlin.
- Reiss, J., Schulze, P., Sesterhenn, J., and Mehrmann, V. (2018). The shifted proper orthogonal decomposition: A mode decomposition for multiple transport phenomena. *SIAM Journal on Scientific Computing*, 40(3):A1322–A1344.
- Silverman, B. W. (1986). *Density Estimation for Statistics and Data Analysis*. Chapman & Hall, London.
- Tsybakov, A. B. (2009). *Introduction to Nonparametric Estimation*. Springer, New York.
- Vicsek, T., Czirók, A., Ben-Jacob, E., Cohen, I., and Shochet, O. (1995). Novel type of phase transition in a system of self-driven particles. *Physical Review Letters*, 75(6):1226–1229.